

About PGA™ — Prediction Governance Architecture

Governing Predictions Across the Complete Prediction Lifecycle

Predictions increasingly influence real-world decisions.

Human forecasters estimate future outcomes.

AI systems generate probabilities, forecasts, risk estimates, and predictive signals.

Prediction markets aggregate expectations from multiple participants.

Organizations rely on forecasts when allocating capital, managing risk, planning infrastructure, operating supply chains, and making strategic decisions.

Autonomous agents and decision systems may increasingly consume predictions and act upon them without continuous human interpretation.

Predictive systems may also influence the reality they attempt to predict.

As prediction becomes embedded across human, organizational, AI, market, and autonomous decision environments, the prediction itself requires governance beyond the model, person, market, or mechanism that generated it.

PGA™ — Prediction Governance Architecture was created to provide a stable governance architecture for this expanding Prediction domain.

PGA™ does not predict the future, prescribe forecasting models, determine Operational Reality, validate evidence, allocate accountability, or replace the systems that generate predictions.

PGA™ governs predictions. It does not make predictions.

Why PGA™ Exists

Predictive systems produce enormous amounts of forecasts, probabilities, estimates, signals, scenarios, and expectations.

But the existence of a prediction does not itself answer fundamental governance questions:

What exactly is being predicted?

For what purpose?

Over what horizon and within what boundary?

By what methodology was the prediction constructed?

What information and signals formed its Prediction Basis?

What assumptions must hold?

How was the specific prediction formed?

What does its probability actually mean?

How uncertain is the prediction?

Is the predicted outcome operationally feasible?

If multiple predictors are combined, are they genuinely diverse and sufficiently independent?

How may the prediction be communicated and relied upon?



When must changing conditions trigger review or revision?

What actually happened, and is that outcome legitimately comparable with the prediction?

Does observed accuracy demonstrate predictive skill?

Can the complete prediction history later be reconstructed?

PGA™ addresses these questions by defining the governable spaces that exist throughout the complete Prediction Lifecycle.

Rather than standardizing one forecasting technology, PGA™ standardizes the governance conditions under which predictions can be defined, methodologically constructed, basis-grounded, assumption-visible, formed, uncertainty- bounded, feasibility-aware, collectively governed where applicable, communicated, relied upon, revised, performance- assessed, and historically reconstructed.

The Prediction Before the Decision

The objective of PGA™ is not to prescribe how an organization, human, AI system, market, or autonomous agent must predict the future.

The objective is to define how the prediction itself must remain governable before, during, and after reliance.

Every prediction requires a defined Prediction Object and purpose.

Every prediction requires a methodology.

Every methodology operates upon a Prediction Basis.

Predictions depend upon assumptions.

Those elements must be transformed into a specific Governed Prediction.

Probability and uncertainty must be represented without creating false certainty.

The predicted outcome may require feasibility assessment.

Multiple predictions may require Collective Prediction Governance before they can legitimately be represented as a collective signal.

Predictions must remain bounded when communicated and relied upon.

Changing conditions may require review, revision, supersession, or withdrawal.

A sufficiently governed outcome may eventually become available.

Predictive Performance may then become assessable.

The complete Prediction Record must remain historically reconstructable.

PGA™ therefore treats the prediction as an independent governance domain surrounding the mechanisms that generate and consume predictive information.

The Prediction Governance Architecture

PGA™ consists of twelve independent Prediction Governance spaces.

Each space answers one architectural question.

Each space governs one distinct part of the Prediction Lifecycle.

Each space has a defined beginning, operational purpose, and governance boundary.

Each space begins where the preceding Prediction Governance requirement has been sufficiently established.

Together, these spaces form the complete Prediction Governance Architecture.

The architecture progresses through:

Prediction Object & Purpose → Prediction Methodology → Prediction Basis → Prediction Assumptions → Prediction Formation → Probability & Uncertainty → Prediction Feasibility → Collective Prediction where applicable → Communication & Reliance → Monitoring & Revision → Outcome & Performance → Prediction Record & Continuity

This progression represents the governance journey from defining what is being predicted and why to determining how that prediction remains methodologically interpretable, appropriately relied upon, performance-assessable, and historically reconstructable throughout time .

Prediction Spaces Rather Than Forecasting- System Components

PGA™ does not organize Prediction Governance around AI models, forecasting platforms, prediction markets, data providers, algorithms, departments, vendors, statistical techniques, or technical components.

Instead, it organizes governance around distinct Prediction Governance spaces .

Each space represents an independent governance responsibility that exists regardless of the predictive mechanism producing the prediction.

A human forecast, AI-generated probability, financial forecast, prediction-market signal, industrial failure prediction, geopolitical forecast, autonomous-agent prediction, or collective forecast may be technically very different.

Yet each may still require governance of:

object & purpose → methodology → basis → assumptions → formation → probability & uncertainty → feasibility → collective prediction where applicable → communication & reliance → monitoring & revision → outcome & performance → record & continuity.

Because these spaces are methodology-neutral and technology-neutral, PGA™ can remain applicable while predictive models, AI systems, markets, forecasting methods, and autonomous technologies continue to evolve.

A Methodology Rather Than a Prediction Engine

PGA™ is a governance methodology.

It is not a prediction engine.

The architecture defines which Prediction Governance conditions must exist, how they relate to one another, where their boundaries lie, and what must remain explicit for a prediction to remain governable throughout its lifecycle .

Organizations and predictive systems remain free to select their own:

forecasting models, statistical methods, machine-learning systems, foundation models, AI agents, simulation systems, prediction markets, data sources, forecasting platforms, probabilistic methods, expert-judgment processes, collective-intelligence mechanisms.

PGA™ governs the methodological Prediction structure surrounding those implementations rather than replacing them.

This separation allows the architecture to remain stable while prediction technologies and methodologies change.

From Predictive Claim to Governed Prediction

A prediction cannot be governed merely as a number, sentence, probability, model output, or market price.

Before prediction formation, the Prediction Object and purpose must be defined.

Before methodology can be interpreted, the Prediction Basis and assumptions must remain explicit.

Before a probability can be relied upon, its uncertainty, confidence, ranges, and conditionality must be understood.

Before predicted outcomes are treated as realistic, feasibility must be considered.

Before multiple predictions become a collective prediction, diversity, independence, concentration, and aggregation conditions must be governed.

Before a prediction influences decisions, communication and reliance boundaries must remain explicit.

Before historical accuracy becomes a performance claim, the prediction must be legitimately associated with a sufficiently governed outcome and interpreted within its original informational and operational conditions.

And before long-term evaluation is possible, revisions, supersessions, withdrawals, outcomes, performance states, and historical versions must remain preserved.

PGA™ transforms these dependencies into a structured and navigable Prediction Lifecycle.

Probability Is Not Certainty

One of the most important distinctions within PGA™ is:

Probability ≠ Certainty

A prediction may state a high probability.

That probability may still contain substantial uncertainty.

Its confidence may be limited.

Its methodology may depend upon fragile assumptions.

Its Prediction Basis may change.

The predicted outcome may even face operational feasibility constraints.

PGA™ therefore preserves the distinctions between:

Probability

Confidence

Uncertainty

Feasibility

and

Certainty.

High Probability ≠ High Confidence.

Same Probability Value ≠ Same Predictive Meaning.

Probability becomes meaningful only within the methodological and informational conditions from which it was derived.

Feasibility Is Not Probability

PGA™ also establishes another fundamental distinction:

Feasible ≠ Probable

An outcome may be possible but operationally unrealistic.

An outcome may be feasible but unlikely.

A highly probable forecast may depend upon assumptions that make the predicted outcome operationally difficult or impossible.

A low-probability outcome may nevertheless remain entirely feasible.

PGA™ therefore distinguishes:

Possible ≠ Plausible ≠ Feasible ≠ Probable

Prediction Feasibility Governance asks whether the predicted outcome can realistically occur under the stated constraints, assumptions, horizon, Prediction Basis, and relevant operational conditions.

It does not replace probability estimation.

Collective Prediction Is Not Automatically Collective Intelligence

PGA™ deliberately separates multiple predictors from legitimate Collective Prediction.

Multiple Predictors ≠ Independent Predictions.

A thousand participants may rely upon the same information source.

Multiple models may share the same training data.

Different forecasters may follow the same dominant narrative.

Prediction-market participants may react to the same information.

Several AI agents may ultimately depend upon the same underlying model.

Large participation therefore does not automatically create information diversity or predictive independence.

Number of Participants ≠ Information Diversity.

A collective signal is not automatically a governed prediction.

PGA™ governs whether multiple predictions may legitimately be combined and what the resulting collective representation actually establishes.

The Prediction-to-Reliance Boundary

The transition from Governed Prediction to real-world reliance is one of the most consequential points in the architecture.

A prediction may be methodologically legitimate while still being inappropriate for a particular decision.

A probability may be communicated accurately while being interpreted incorrectly.

A forecast may be suitable for strategic planning but unsuitable for immediate operational action.

A prediction may become outdated while remaining publicly visible.

These conditions matter.

But the governance question is not simply:

What does the prediction say?

The additional question is:

What does the prediction legitimately support?

PGA™ therefore treats the transition from Prediction Communication to Prediction Reliance as an explicit governance boundary.

A prediction must not become stronger in communication than it is in methodology.

Communicable ≠ Reliance-Eligible.

Lifecycle Governance and Prediction Continuity

Prediction Governance does not end when a prediction is published.

Information changes.

Signals change.

Assumptions may fail.

Operational constraints may emerge.

Probability may change.

Feasibility may change.

A prediction may influence behavior and thereby influence Operational Reality.

Review may become necessary.

A prediction may be revised.

A later prediction may supersede an earlier state.

A prediction may be withdrawn.

An outcome may become available.

Performance may become assessable.

A historical prediction may later become relevant again.

PGA™ therefore extends governance beyond Prediction Formation into Prediction Monitoring & Revision, Prediction Outcome & Performance, and Prediction Record & Continuity .

Where material change affects the prediction, the architecture preserves the earlier Prediction State rather than rewriting it.

Where a historical prediction later returns to use, PGA™ supports governed lifecycle re-entry.

The result is not uncontrolled reconstruction of the prediction.

It is targeted governed Prediction lifecycle re-entry .

Who May Benefit

PGA™ may support:

AI developers, AI operators, predictive-model providers, forecasting organizations, prediction platforms, prediction markets, financial institutions, investment organizations, insurers, industrial organizations, infrastructure operators, supply-chain organizations, risk-management teams, strategic-planning teams, public institutions, research organizations, autonomous-system developers, agentic-system operators, decision-support providers, collective-intelligence platforms, audit and assurance organizations, organizations relying upon predictions for material decisions.

Relationship to Other OOF® Architectures

PGA™ governs predictions.

It operates alongside complementary OOF® architectures.

GOA™ — Governance Architecture governs governance structure, authority, scope, delegation, accountability, continuity, and interoperability.

ORA™ — Operational Reality Architecture governs Operational Reality and provides the methodological distinction between what actually occurs and what was previously predicted.

VALIDOS™ — Validation Governance Architecture governs validation purpose, criteria, methods, evidence fitness, determination, Validation States, reliance, and revalidation.

INTEGROS® — Integrity Governance Architecture governs integrity across the complete integrity lifecycle.

ADIT® — Audit in Real Time Architecture governs evidential observation, preservation, verification, examination, traceability, and auditability.

AGA™ — Accountability Governance Architecture governs accountability, responsibility, attribution, delegation, and accountable relationships.

AIG® — AI Governance Architecture governs AI behavior.

ASGA™ — Autonomous Systems Governance Architecture governs autonomous systems, autonomous agents, authority, constraints, execution, and operational autonomy.

CLIA® governs cognition-related governance spaces relevant to intelligent systems.

SIMULOS® — Simulation Governance Architecture governs simulation purpose, structure, execution, outputs, and lifecycle.

Together, these architectures allow Prediction Governance to interact with Operational Reality, validation, integrity, evidence, accountability, AI behavior, autonomy, cognition, and simulation without collapsing those independent governance domains into one architecture.

Use Case 1 — Prediction Market and Information Advantage

A prediction market displays an 86% probability that a major corporate transaction will close before a specified date.

Thousands of participants have traded.

The market is widely described as demonstrating the predictive intelligence of the crowd.

Several highly profitable participants appear to have exceptional historical prediction records.

PGA™ asks an earlier set of questions.

What exactly is being predicted?

What does the 86% value methodologically represent?

Which information forms the Prediction Basis?

How independent are the participating predictions?

Do participants rely upon common information sources?

Is market participation representative of broader belief?

How concentrated is predictive influence?

How close are the highest-performing predictors to the actors, information, decisions, or operational processes from which the outcome may emerge?

Did some predictors possess information unavailable to the broader market?

Did any predictor know the likely outcome rather than predict it through superior predictive methodology?

Was the 86% probability stable, or did it converge only shortly before resolution?

PGA™ preserves the distinctions:

Market-Implied Probability ≠ Governed Probability Estimate

Market-Implied Probability ≠ Population Belief

Information Advantage $\neq$ Predictive Capability

Outcome Knowledge $\neq$ Predictive Capability

Observed Accuracy $\neq$ Demonstrated Predictive Skill

The objective is not to accuse information-advantaged participants of misconduct.

High Predictor–Reality Proximity $\neq$ Misconduct.

The objective is to ensure that historical predictive performance is interpreted within the informational conditions under which the prediction actually existed.

Use Case 2 — AI Prediction That Changes the Outcome

An industrial AI system predicts a 75% probability of critical equipment failure within 30 days .

The organization relies upon the prediction.

Maintenance teams immediately inspect the equipment, replace a deteriorating component, and prevent the failure.

Thirty days later, no failure occurs.

A simplistic performance system records:

Prediction: Failure

Outcome: No Failure

Prediction Result: Incorrect

PGA™ asks a different set of questions.

What was the Prediction Basis when the prediction was formed?

What assumptions supported the 75% probability?

Was failure operationally feasible under the original conditions?

When was the prediction communicated?

What action did reliance upon the prediction trigger?

Did that intervention materially alter Operational Reality?

What outcome would have remained plausible without intervention?

Should the observed absence of failure be interpreted as prediction error, intervention-affected outcome, or another qualified performance state?

The prediction may have contributed to preventing the outcome it predicted.

PGA™ therefore preserves the distinctions:

Prediction Observation $\neq$ Prediction Influence

Prediction Influence $\neq$ Outcome Control

Prediction Influence $\neq$ Predictive Accuracy

Predictive Accuracy $\neq$ Operational Value

The absence of the predicted failure does not automatically establish that the prediction was methodologically poor.

A prediction capable of materially influencing the reality it predicts must govern that influence as part of its continuing predictive conditions.

The Long-Term Vision

Predictive activity is moving from periodic human forecasts toward continuously operating AI, agentic, market, collective, and autonomous predictive environments.

AI systems may generate millions of predictions.

Agents may consume predictions from other agents.

Prediction markets may continuously aggregate changing expectations.

Autonomous systems may predict operational conditions and act before humans review the forecast.

Models may update probabilities continuously as new information arrives.

Predictions may influence the behavior that determines whether those predictions become true.

One predictive system may generate thousands of changing Prediction States before an outcome becomes observable.

One final prediction may hide months or years of earlier predictive history.

As this transition accelerates, the central question will no longer be only whether systems can generate more accurate predictions.

It will also be:

How can we know what was actually predicted, by what methodology, from what information, under what assumptions, with what uncertainty, before the outcome became known?

The long-term objective of PGA™ is to provide a common Prediction Governance language through which humans, organizations, AI systems, markets, autonomous agents, public institutions, and future predictive infrastructures can define, compare, govern, communicate, evaluate, reconstruct, and exchange predictions without confusing probability with certainty, information advantage with predictive capability, observed accuracy with predictive skill, or prediction with Operational Reality.

PGA™ is intended to make predictions governable before reliance, during changing predictive conditions, after outcomes become known, throughout performance assessment, and across the complete historical Prediction Lifecycle.

Canonical Principle

Predictive models, AI systems, prediction markets, autonomous agents, forecasting methodologies, information environments, and decision technologies will continue to change. Prediction Governance must remain stable. PGA™ therefore governs the methodological conditions that determine what is being predicted and why, how the prediction is constructed, what information and assumptions support it, how probability and uncertainty are represented, whether the predicted outcome is feasible, when multiple predictions may legitimately become a collective prediction, how predictions may be communicated and relied upon, when changing conditions require review or revision, how predictions are associated with governed outcomes and assessed for Predictive Performance, and how their complete historical record remains reconstructable throughout the Prediction Lifecycle.

Canonical Closing Statement

PGA™ — Prediction Governance Architecture provides the governance foundation through which human, organizational, model-generated, AI-generated, market-derived, collective, agentic, and autonomous predictions can be explicitly defined, methodologically documented, grounded in a governed Prediction Basis, associated with visible assumptions, formed into Governed Predictions, represented through bounded probability and uncertainty, assessed for feasibility, collectively governed where applicable, communicated and relied upon within legitimate boundaries, monitored and revised as predictive conditions change, associated with sufficiently governed outcomes, evaluated for Predictive Performance, and continuously preserved through complete Prediction Record & Historical Continuity. By defining the independent governable spaces of the complete

Prediction Lifecycle, PGA™ establishes a methodology-neutral and technology-neutral architecture for ensuring that predictions remain interpretable without requiring one forecasting method, model, market, or predictive technology to define Prediction Governance itself. PGA™ therefore establishes Prediction Governance as a distinct methodological layer around the prediction, enabling predictions to remain traceable, uncertainty- bounded, reliance-aware, performance-assessable, and historically reconstructable regardless of whether they originate from humans, AI systems, markets, autonomous agents, or future predictive mechanisms.

PGA™ governs the prediction — not the future.

PGA™ Parent Standards

[Prediction Object & Purpose Standard \(POPS\)](#)

[Prediction Methodology Standard \(PMES\)](#)

[Prediction Basis Standard \(PBAS\)](#)

[Prediction Assumption Standard \(PAS\)](#)

[Prediction Formation Standard \(PFOS\)](#)

[Prediction Probability & Uncertainty Standard \(PPUS\)](#)

[Prediction Feasibility Standard \(PFES\)](#)

[Collective Prediction Governance Standard \(CPGS\)](#)

[Prediction Communication & Reliance Standard \(PCRS\)](#)

[Prediction Monitoring & Revision Standard \(PMRS\)](#)

[Prediction Outcome & Performance Standard \(POFPS\)](#)

[Prediction Record & Continuity Standard \(PRCS\)](#)

1

NON-EXECUTABLE NOTICE

OOF® — OriginOpen® Foundation is a non-executable methodological authority.
This website and all publications are reference-only.

CANONICAL LANGUAGE NOTICE

UCL™ is the canonical language layer of OOF® publications.
English is its carrier language.
In case of ambiguity, the published English HTML text prevails.

RIGHTS & PROTECTION

© OOF® — OriginOpen® Foundation
Protected under MIP® — Methodological Intellectual Property.
All rights reserved.



contact@originopenfoundation.org