

## PCALM — Prediction Calibration Module

Architecture Ecosystem: Structured Reality Standards™

Architecture Family: PGA™ — Prediction Governance Architecture

### Parent Standard:

POFPS — Prediction Outcome & Performance Standard

Operational Layer: Prediction Outcome & Predictive Performance

### Category:

Governance & Enforcement

### Subcategory:

Prediction Governance

Type: Parent Standard Module

### Governed Space:

Prediction Calibration

Version: 1.0

Status: Canonical · Open Module

Effective Date: 11 September 2026

Compatibility: OOF® Methodology OS™ · GOA™ · ORA™ · VALIDOS™ · INTEGROS® · ADIT® · AGA™ · AIG® · ASGA™ · CLIA® · SIMULOS®

AI-Readable: Yes

Authority: OOF®

### Protection:

MIP® — Methodological Intellectual Property

### Canonical Language:

English (UCL™)

## Governed Space

PCALM governs how probabilistic predictions are grouped, compared with sufficiently governed outcomes, assessed for correspondence between stated probability and observed frequency, segmented, qualified, versioned, and preserved as a legitimate Calibration Record.

It governs the transition:

Governed Probabilistic Prediction Population + Governed Outcome Associations + Calibration Method → Governed Prediction Calibration State

PCALM does not ask whether one prediction was correct.

It asks whether repeated predictions assigned comparable probabilities resolve at frequencies consistent with what those probabilities represented.

Its foundational principle is:

A probability becomes calibration-relevant only across a governed population of comparable predictions.

## Canonical Definition

The Prediction Calibration Module (PCALM) defines the operational governance requirements by which probabilistic predictions are assembled into legitimate Calibration Populations and assessed to determine whether stated probabilities correspond, over time and across sufficiently governed cases, to observed outcome frequencies without allowing selective inclusion, hidden segmentation, unstable bins, post-outcome regrouping, incomparable prediction types, or outcome leakage to manufacture apparent calibration.

PCALM governs:

Calibration Eligibility; Calibration Population;

Calibration Unit; Calibration Window; probability grouping; probability bands; calibration bins; bin stability; observed frequency; expected frequency; calibration gap; overconfidence; underconfidence; local calibration; global calibration; segmented calibration; horizon-specific calibration; domain-specific calibration; predictor-specific calibration; methodology-specific calibration; model-specific calibration; collective-prediction calibration; prediction-market calibration; sample sufficiency; independence context; weighting; calibration drift; recalibration; calibration history; calibration qualification.

The primary output is:

Governed Prediction Calibration Record

## Operational Role

PCALM is the fourth operational module within POFPS.

Its focused operational question is:

Across a sufficiently governed population of probabilistic predictions, do stated probability levels correspond to observed outcome frequencies in a way that legitimately supports calibration claims?

PCALM receives:

governed historical probabilistic predictions; POAM-governed Outcome Associations; POCM comparison states; PEMM error or event-level probability-loss information where relevant; historical version context.

It produces calibration information for:

PPCM — Predictive Performance Module

## Internal Module Flow

Prediction Outcome Association → Prediction Outcome Comparison → Prediction Error Measurement → Prediction Calibration → Predictive Performance

### Current Position:

Prediction Calibration

## 1. Prediction Calibration

Prediction Calibration concerns the relationship between:

Stated Probability

and:

Observed Outcome Frequency

across a sufficiently governed set of comparable probabilistic predictions.

Conceptually:

If many predictions are assigned approximately:

70%

then approximately the expected proportion should eventually resolve positively over a sufficiently meaningful population.

## 2. Calibration Eligibility

Not every probabilistic statement is eligible for calibration.

A prediction may be:

Calibration-Eligible

Calibration-Eligible with Qualification

Partially Eligible

Not Yet Eligible

Not Comparable for Calibration

Unable to Determine

## 3. Calibration Requires Repetition

One prediction cannot establish calibration.

Example:

70% probability

followed by occurrence or non-occurrence says almost nothing about whether the predictor is calibrated.

Single Probability Error  $\neq$  Calibration

## 4. Calibration Population

A Calibration Population is the governed set of predictions used to assess probability-frequency correspondence.

It should preserve:

prediction identities; prediction versions;

domains; horizons; methodologies; probability values; outcome references; inclusion and exclusion rules.

## 5. Population Definition Before Analysis

Where possible, the Calibration Population should be defined before final performance results are inspected.

This reduces post-outcome selection.

## 6. Population Stability

A Calibration Population should not be silently redefined after outcomes are known.

Post-Outcome Population Change  $\neq$  Neutral Calibration Update

## 7. Calibration Unit

The Calibration Unit identifies the atomic predictive case being counted.

Examples:

one event prediction; one prediction version; one final prediction per event; one original prediction per event; one decision-time prediction per event.

The unit must remain explicit.

## 8. Prediction Version Rule

Where predictions are revised, calibration must specify whether the population contains:

original versions; final versions; every material version; decision-time versions; another governed state.

Final-Version Calibration  $\neq$  Original-Version Calibration

## 9. Event Duplication

Repeated updates concerning the same event can inflate sample size.

PCALM should identify whether multiple prediction versions are being counted as independent cases.

Prediction Version Count  $\neq$  Independent Event Count

## 10. Calibration Window

Calibration may be assessed over a defined period.

Examples:

Quarterly

Annual

Rolling 12-Month

Model-Version Window

Methodology-Version Window

The window should remain explicit.

## 11. Rolling Calibration

Rolling windows may reveal changing calibration over time.

But overlapping windows can create dependencies that should remain visible.

## 12. Probability Band

Probabilities may be grouped into bands such as:

0–10%

10–20%

and so on.

The banding method should be governed.

## 13. Calibration Bin

A Calibration Bin is a governed grouping of predictions with sufficiently similar stated probabilities.

Example:

65–75%

may be used to approximate a 70% calibration band.

## 14. Bin Width

Very wide bins can hide local miscalibration.

Very narrow bins can produce insufficient sample size.

PCALM should balance:

Resolution

and:

Sample Sufficiency

## 15. Fixed Bins

Fixed bins use predetermined probability intervals.

They support stability and comparability across time.

## 16. Adaptive Bins

Adaptive bins may change based on available prediction density.

They can improve sample size but require strong traceability.

## 17. Post-Outcome Bin Manipulation

Bins should not be redesigned after outcomes are known simply to produce better-looking calibration.

## 18. Expected Frequency

For a group of predictions averaging approximately:

70%

the expected positive-outcome frequency is approximately:

70%

subject to the defined grouping method.

## 19. Observed Frequency

Observed Frequency is the governed proportion of positive outcomes within the Calibration Population or Calibration Bin.

## 20. Calibration Gap

Calibration Gap represents the difference between:

Stated / Mean Probability

and:

Observed Frequency

within the governed population or bin.

## 21. Positive Calibration Gap

Depending on sign convention, a positive gap may indicate overprediction of event frequency.

The convention must remain explicit.

## 22. Negative Calibration Gap

A negative gap may indicate underprediction.

Again, sign convention must be documented.

## 23. Absolute Calibration Gap

Absolute Calibration Gap removes direction and preserves magnitude.

Direction should still remain available separately.

## 24. Perfect Calibration Boundary

A calibration gap of zero within one bin or one period does not prove perfect calibration generally.

Local Zero Gap  $\neq$  Universal Calibration

## 25. Overconfidence

Overconfidence may exist where stated probabilities are systematically more extreme than observed outcome frequencies justify.

Example:

Predictions near:

90%

resolve positively only:

65%

of the time.

## 26. Underconfidence

Underconfidence may exist where stated probabilities are systematically too conservative relative to observed frequency.

## 27. Calibration Direction

Calibration Direction may indicate:

Overconfident

Underconfident

Approximately Calibrated

Mixed

Unable to Determine

## 28. Local Calibration

Local Calibration assesses specific probability regions.

Example:

A predictor may be well calibrated at:

40–60%

but poorly calibrated at:

80–95%

## 29. Global Calibration

Global Calibration summarizes calibration across the full governed probability range.

Global Calibration  $\neq$  Local Calibration in Every Probability Region

### **30. Calibration Curve**

A Calibration Curve may represent:

Predicted Probability  $\rightarrow$  Observed Frequency  
across bins.

The curve is a representation of governed calibration data.

It is not itself a proof of predictive skill.

### **31. Calibration Slope**

Some implementations may estimate Calibration Slope.

The exact statistical method is implementation-specific and must remain documented.

### **32. Calibration Intercept**

Calibration Intercept may also be used where appropriate.

Again, PCALM governs methodological transparency rather than prescribing one universal statistic.

### **33. Calibration Error Metric**

Implementations may use governed metrics such as:

mean absolute calibration error; weighted calibration error; expected calibration error; maximum calibration gap; other appropriate metrics.

No universal single metric is mandated.

### **34. Metric Definition**

Every Calibration Metric should identify:

population; bins; weighting; calculation; sample size; interpretation limits.

Calibration Score  $\neq$  Self-Explanatory Calibration

### **35. Weighted Calibration**

Calibration bins may be weighted by:

number of predictions; event importance; another governed basis.

The weighting must remain explicit.

### **36. Equal-Bin Weighting**

Equal weighting of bins may give sparse bins disproportionate influence.

This should remain visible.

### **37. Prediction-Count Weighting**

Weighting by prediction count emphasizes where most predictions occur.

It may underrepresent calibration in rare probability regions.

### **38. Consequence Weighting Boundary**

High-consequence predictions may deserve separate analysis.

But consequence weighting changes the question being answered.

Calibration Frequency  $\neq$  Consequence-Weighted Risk Performance

### **39. Sample Sufficiency**

Calibration claims require enough observations to support the proposed conclusion.

A bin containing:

3 predictions

cannot support the same confidence as one containing:

3,000 predictions.

### **40. Sparse Bin**

A Sparse Bin has insufficient observations for strong calibration inference.

Possible state:

Calibration Indeterminate — Sparse Sample

### **41. Sample Size Disclosure**

Every material Calibration Result should preserve sample size.

### **42. Small-Sample Volatility**

Observed frequency can vary substantially in small samples.

PCALM should prevent small-sample noise from being interpreted as stable calibration.

### **43. Confidence Interval Context**

Implementations may preserve statistical intervals or uncertainty around observed frequency where appropriate.

PCALM does not mandate one statistical method.

### **44. Calibration Uncertainty**

Calibration itself can be uncertain.

Therefore:

Calibration Estimate  $\neq$  Perfectly Known Calibration

### **45. Population Comparability**

Predictions should be sufficiently comparable before being grouped.

Relevant dimensions may include:

Prediction Object; domain; horizon; methodology; environment; population; probability representation.

### **46. Cross-Domain Calibration**

Combining unrelated domains may conceal domain-specific calibration.

Example:

Good weather calibration may hide poor technology-forecast calibration.



## 47. Domain-Specific Calibration

PCALM may produce separate calibration states by domain.

## 48. Horizon-Specific Calibration

Predictors may be well calibrated at short horizons and poorly calibrated at long horizons.

Aggregate Calibration  $\neq$  Horizon-Specific Calibration

## 49. Methodology-Specific Calibration

Predictions generated under different methodologies should remain separately identifiable where differences materially affect calibration.

## 50. Model-Specific Calibration

AI model versions may have distinct calibration characteristics.

Model Upgrade  $\neq$  Calibration History Reset

## 51. Predictor-Specific Calibration

Human or machine predictors may have separately governed calibration records.

## 52. Team Calibration

A forecasting team may have group-level calibration distinct from individual members.

## 53. Collective Prediction Calibration

Collective predictions may be calibrated against outcomes.

However, CPGS context should remain visible.

## 54. Market-Implied Probability Calibration

Prediction-market probabilities may be assessed for calibration where:

market state is preserved; timing is explicit; resolution criteria are governed; selected prediction points are defined.

## 55. Market Snapshot Rule

Calibration should specify which market probability is used.

Examples:

opening probability; daily close; decision-time probability; final pre-resolution probability.

Final Market Probability  $\neq$  Entire Market Prediction History

## 56. Time-to-Outcome Calibration

Calibration may vary by time remaining before resolution.

A predictor may be poorly calibrated six months out but strongly calibrated one day out.

## 57. Lead-Time Segmentation

PCALM may segment predictions by lead time.

This prevents late-stage information advantage from obscuring earlier predictive quality.

## 58. Revision-Aware Calibration

Where predictions are revised, calibration should specify the version rule.

Possible calibration views include:

Original-Version Calibration

Final-Version Calibration

Decision-Time Calibration

All-Version Calibration

## 59. Late-Revision Effect

Final-version calibration may appear excellent if probabilities are updated immediately before resolution.

This does not automatically demonstrate strong long-horizon calibration.

## 60. Withdrawal-Aware Calibration

Withdrawn predictions should not automatically disappear from the Calibration Population.

Whether they remain calibration-eligible depends on the specific calibration question and available outcomes.

Withdrawal  $\neq$  Automatic Calibration Exclusion

## 61. Supersession-Aware Calibration

Superseded predictions may remain part of historical calibration analysis.

## 62. Selective Inclusion Risk

Removing difficult or unfavorable predictions can improve apparent calibration.

PCALM requires explicit inclusion rules.

## 63. Selective Exclusion Risk

Legitimate exclusions may exist.

They must remain documented and traceable.

## 64. Survivorship Bias

Calibration based only on predictions surviving until resolution may be distorted.

Visible Resolved Predictions  $\neq$  Complete Calibration Population

## 65. Publication Bias

Published predictions may differ systematically from unpublished predictions.

Calibration claims should identify which population is being assessed.

## 66. Missing Outcome Bias

Some predictions may never obtain sufficiently governed outcomes.

Their exclusion may affect calibration interpretation.

The missingness should remain visible.

## 67. Outcome Resolution Bias

Outcomes easier to resolve may be overrepresented in calibration datasets.

## **68. Predictor–Reality Proximity Context**

Calibration may appear strong where predictors are close to outcome-determining information.

PCALM preserves PRP context where material.

## **69. Information-Advantaged Calibration**

A predictor with privileged information may still be genuinely calibrated.

But that does not necessarily prove generalizable predictive capability.

Good Calibration  $\neq$  Information Independence

## **70. Outcome-Informed Calibration**

If predictors already know substantial outcome-determining information, calibration interpretation changes.

## **71. Outcome-Controlled Calibration**

A predictor controlling outcomes may appear extremely calibrated.

Outcome Control  $\neq$  Predictive Skill

Calibration can be numerically real while interpretively limited.

## **72. Self-Fulfilling Calibration Context**

If predictions systematically influence outcomes, observed calibration may partly reflect that influence.

## **73. Self-Negating Calibration Context**

If high-probability predictions trigger interventions that prevent events, raw calibration may appear poor unless intervention context is preserved.

## **74. Intervention-Aware Calibration**

PCALM may segment or qualify predictions where interventions materially altered the environment.

It should not invent counterfactual outcomes.

## **75. Counterfactual Boundary**

Calibration uses observed governed outcomes.

It does not replace them with hypothetical outcomes.

Counterfactual Calibration  $\neq$  Observed Calibration

## **76. Correlated Predictions**

Predictions may share:

same data; same methodology; same event; same model; same predictor.

This reduces independence.

## **77. Independence Context**

PCALM should preserve whether predictions are:

Independent

Partially Dependent

Strongly Dependent

Unable to Determine

where relevant to interpretation.

## 78. Duplicate Event Predictions

Multiple forecasts of the same event should not automatically be treated as independent calibration cases.

## 79. Clustered Predictions

Predictions may cluster around a common underlying event.

Calibration analysis should preserve cluster structure where material.

## 80. Calibration Drift

Calibration Drift occurs where probability-frequency correspondence changes over time.

Possible causes may include:

methodology change; environment change; model change; predictor change; data change; regime change.

PCALM identifies drift.

It does not automatically determine cause.

## 81. Drift Detection

Calibration may be compared across windows to detect:

Improvement

Degradation

Stable Calibration

Mixed Drift

Unable to Determine

## 82. Recalibration

A predictive system may undergo recalibration.

Examples include:

probability transformation; model recalibration; human forecast adjustment; ensemble correction.

Recalibration should create a new governed state.

## 83. Recalibration Boundary

Recalibration changes future probability interpretation.

It does not rewrite historical probabilities.

Recalibration  $\neq$  Historical Probability Replacement

## 84. Calibration Before vs After Recalibration

Performance should preserve:

Pre-Recalibration Calibration

and:

Post-Recalibration Calibration

separately.

## 85. Calibration Benchmark

Calibration may be compared against:

prior predictor version; benchmark model; market; base-rate system; another governed forecaster.

## 86. Benchmark Comparability

A benchmark should use sufficiently comparable prediction populations.

Different Calibration Population  $\neq$  Fair Calibration Benchmark by Default

## 87. Calibration Ranking

Ranking predictors by calibration alone may conceal differences in:

resolution; sharpness; information value; domain difficulty; horizon; sample size.

Best Calibrated  $\neq$  Best Predictor by Default

## 88. Calibration and Resolution

A predictor assigning:

50% to everything

may appear calibrated in a balanced environment while providing little discriminatory value.

Therefore:

Calibration  $\neq$  Resolution

## 89. Calibration and Sharpness

Sharpness concerns how concentrated or decisive probability distributions are.

Strong sharpness without calibration can be dangerous.

Calibration without useful sharpness can be uninformative.

## 90. Calibration and Predictive Skill

Calibration is one dimension of predictive performance.

It is not the whole of predictive skill.

Well-Calibrated  $\neq$  Highly Informative by Default

## 91. Governance Requirements

A governed PCALM implementation should be able to:

1. identify Calibration-Eligible predictions;
2. define a Calibration Population;
3. identify the Calibration Unit;
4. define the Prediction Version Rule;
5. prevent event duplication from inflating sample size;
6. establish Calibration Window;
7. support Rolling Calibration where appropriate;
8. define Probability Bands;
9. define Calibration Bins;
10. preserve Bin Width;
11. support Fixed Bins;
12. support Adaptive Bins;
13. prevent Post-Outcome Bin Manipulation;
14. calculate Expected Frequency;
15. calculate Observed Frequency;
16. calculate Calibration Gap;
17. preserve Calibration Gap direction;
18. calculate Absolute Calibration Gap;
19. prevent Local Zero Gap from being treated as universal calibration;
20. identify Overconfidence;
21. identify Underconfidence;
22. determine Calibration Direction;
23. support Local

Calibration; 24. support Global Calibration; 25. preserve Calibration Curve; 26. support Calibration Slope where appropriate; 27. support Calibration Intercept where appropriate; 28. support governed Calibration Error Metrics; 29. preserve Calibration Metric definition; 30. govern Weighted Calibration; 31. distinguish Equal-Bin from Prediction-Count Weighting; 32. preserve Consequence Weighting Boundary; 33. assess Sample Sufficiency;

34. identify Sparse Bins; 35. disclose Sample Size; 36. preserve Small-Sample Volatility; 37. support statistical uncertainty representation where appropriate; 38. preserve Calibration Uncertainty; 39. assess Population Comparability; 40. preserve Cross-Domain Calibration boundaries; 41. support Domain-Specific Calibration; 42. support Horizon-Specific Calibration; 43. support Methodology-Specific Calibration; 44. support Model-Specific Calibration; 45. support Predictor-Specific Calibration; 46. support Team Calibration; 47. support Collective Prediction Calibration; 48. support Market-Implied Probability Calibration; 49. define Market Snapshot Rule; 50. support Time-to-Outcome Calibration; 51. support Lead-Time Segmentation; 52. support Revision-Aware Calibration; 53. distinguish Original, Final, Decision-Time, and All-Version Calibration; 54. identify Late-Revision Effect; 55. support Withdrawal-Aware Calibration; 56. support Supersession-Aware Calibration; 57. prevent Selective Inclusion; 58. preserve legitimate Selective Exclusion; 59. prevent Survivorship Bias; 60. preserve Publication Bias context; 61. preserve Missing Outcome Bias context; 62. preserve Outcome Resolution Bias context; 63. preserve Predictor-Reality Proximity Context; 64. preserve Information-Advantaged Calibration context; 65. preserve Outcome-Informed Calibration context; 66. preserve Outcome-Controlled Calibration context; 67. preserve Self-Fulfilling Calibration Context; 68. preserve Self-Negating Calibration Context; 69. support Intervention-Aware Calibration; 70. enforce Counterfactual Boundary; 71. identify Correlated Predictions; 72. preserve Independence Context; 73. identify Duplicate Event Predictions; 74. identify Clustered Predictions; 75. detect Calibration Drift; 76. classify Calibration Drift; 77. govern Recalibration; 78. preserve Recalibration Boundary; 79. distinguish Pre- and Post-Recalibration Calibration; 80. support Calibration Benchmarks; 81. assess Benchmark Comparability; 82. prevent simplistic Calibration Ranking; 83. distinguish Calibration from Resolution; 84. distinguish Calibration from Sharpness; 85. distinguish Calibration from Predictive Skill; 86. preserve calibration methodology; 87. preserve calibration qualifications; 88. preserve calibration versioning; 89. preserve Calibration History; 90. preserve calibration traceability; 91. produce a Governed Prediction Calibration Record.

## 92. Governance Logic

PCALM applies the following operational logic:

Is the prediction probabilistic and sufficiently outcome-associated?

If no:

→ calibration does not proceed.

If yes:

→ determine Calibration Eligibility.

What is the Calibration Population?

→ define prediction population and inclusion rules.

What is the Calibration Unit?

→ original prediction, final version, decision-time version, every governed version, or another explicit unit.

Are the predictions sufficiently comparable?

If no:

→ segment or restrict calibration.

What Calibration Window applies?

→ define temporal scope.

How are probabilities grouped?

→ establish bins or another governed calibration structure.

Is sample size sufficient?

If no:

→ classify:

Calibration Indeterminate — Sparse Sample

or equivalent qualified state.

If yes:

→ calculate Observed Frequency.

How does Observed Frequency compare with Stated Probability?

→ determine Calibration Gap.

Is the system systematically overconfident or underconfident?

→ preserve direction and magnitude.

Does calibration vary by domain, horizon, predictor, methodology, model, or lead time?

If yes:

→ segment calibration.

Do revision, withdrawal, supersession, intervention, PRP, or outcome-control conditions materially affect interpretation?

If yes:

→ preserve qualification.

Has calibration materially changed over time?

If yes:

→ identify Calibration Drift.

Then:

→ preserve Calibration Population;

→ preserve method and bins;

→ preserve Sample Size;

→ preserve observed frequency;

→ preserve Calibration Gap;

→ preserve qualifications;

→ preserve history;

→ produce Governed Prediction Calibration Record.

Canonical transition:

Governed Probabilistic Prediction Population + Governed Outcomes → Probability-Frequency Comparison

→ Governed Prediction Calibration Record

## 93. Minimum Implementation Framework

A. Calibration Population

Governed set of predictions included in the calibration assessment.

B. Calibration Unit

Atomic prediction state counted within the population.

C. Prediction Version Rule

Original, Final, Decision-Time, All-Version, or another governed rule.

D. Calibration Window

Time period over which calibration is assessed.

E. Probability Group / Bin

Defined probability interval or grouping.

F. Prediction Count

Number of governed predictions in the applicable bin or population.

G. Expected Frequency

Frequency implied by the stated probabilities.

H. Observed Frequency

Governed realized outcome frequency.

I. Calibration Gap

Difference between expected and observed frequency.

J. Calibration Direction

Approximately Calibrated, Overconfident, Underconfident, Mixed, or Unable to Determine.

K. Sample Sufficiency

Assessment of whether available cases support the proposed calibration conclusion.

L. Segmentation Profile

Domain, horizon, predictor, methodology, model, lead-time, or other material segmentation.

M. Calibration Qualification

Sparse sample, intervention, PRP, outcome control, dependence, missing outcomes, or other limitations.

N. Calibration State

Calibrated, Approximately Calibrated, Miscalibrated, Mixed, Indeterminate, or Unable to Determine, subject to the applicable methodology.

O. Calibration History

Versioned historical record of calibration assessments, drift, and recalibration.

## 94. Governance Outputs

PCALM may produce:

Governed Prediction Calibration Record Calibration Eligibility State Calibration Population Calibration Unit Prediction Version Rule Calibration Window Probability Band Structure Calibration Bin Structure Expected Frequency Observed Frequency Calibration Gap Absolute Calibration Gap Calibration Direction Overconfidence Indicator Underconfidence Indicator Local Calibration Profile Global Calibration Profile Calibration Curve Calibration Error Metric Sample Sufficiency State Sparse Bin Indicator Domain-Specific



Calibration Horizon-Specific Calibration Predictor-Specific Calibration Methodology-Specific Calibration Model-Specific Calibration

Collective Prediction Calibration Market-Implied Probability Calibration Lead-Time Calibration Revision-Aware Calibration Record Withdrawal-Aware Calibration Record Supersession-Aware Calibration Record Intervention-Aware Calibration Record Predictor-Reality Proximity Calibration Context Independence Context Calibration Drift State Recalibration Record Calibration Benchmark Record Calibration Qualification Calibration History Calibration Traceability Record

The primary output is:

Governed Prediction Calibration Record

This becomes a governed input to:

PPCM — Predictive Performance Module

## 95. Critical Distinctions

Single Probability Error  $\neq$  Calibration

Calibration  $\neq$  Single-Prediction Correctness

A Probability Value  $\neq$  Calibration Evidence

Final-Version Calibration  $\neq$  Original-Version Calibration

Prediction Version Count  $\neq$  Independent Event Count

Post-Outcome Population Change  $\neq$  Neutral Calibration Update

Local Zero Gap  $\neq$  Universal Calibration

Global Calibration  $\neq$  Local Calibration in Every Probability Region

Calibration Score  $\neq$  Self-Explanatory Calibration

Calibration Frequency  $\neq$  Consequence-Weighted Risk Performance

Calibration Estimate  $\neq$  Perfectly Known Calibration

Aggregate Calibration  $\neq$  Horizon-Specific Calibration

Model Upgrade  $\neq$  Calibration History Reset

Final Market Probability  $\neq$  Entire Market Prediction History

Withdrawal  $\neq$  Automatic Calibration Exclusion

Visible Resolved Predictions  $\neq$  Complete Calibration Population

Good Calibration  $\neq$  Information Independence

Outcome Control  $\neq$  Predictive Skill

Counterfactual Calibration  $\neq$  Observed Calibration

Recalibration  $\neq$  Historical Probability Replacement

Different Calibration Population  $\neq$  Fair Calibration Benchmark by Default

Best Calibrated  $\neq$  Best Predictor by Default

Calibration  $\neq$  Resolution

Calibration  $\neq$  Sharpness

Calibration  $\neq$  Predictive Skill

Well-Calibrated  $\neq$  Highly Informative by Default

## 96. Enterprise Application

PCALM can govern probability calibration across:

AI forecasting systems; financial forecasting; investment analysis; credit and risk forecasting; insurance; predictive maintenance; demand forecasting; supply chains; logistics; energy systems; technology forecasting; public policy; election forecasting; healthcare operations; prediction markets; model ensembles; Human-AI forecasting; expert forecasting teams; autonomous agents; multi-agent predictive systems.

Its function remains constant:

determine whether stated probabilities behave like the probabilities they claim to be across a sufficiently governed prediction population.

## 97. Use Case 1 — An AI System Is “90% Confident” but Its 90% Predictions Resolve Only 62% of the Time

An enterprise AI forecasting system produces thousands of event predictions.

For predictions assigned between:

88% and 92%

the organization reports:

“These are high-confidence predictions.”

PCALM constructs a governed Calibration Population:

Probability Band: 88–92%

Prediction Count: 640

Outcome-Resolved Cases: 618

Observed Positive Outcome Frequency: 62%

The critical governance question is:

Can a system legitimately represent these predictions as approximately 90% probabilities if comparable events occur only around 62% of the time?

PCALM determines a material Calibration Gap.

The state is:

Overconfidence / Material Miscalibration

subject to the governed calculation method and sampling qualification.

This does not automatically prove:

the model is useless; the methodology is invalid; every 90% prediction is wrong; no operational value exists.

It establishes something narrower and more important:

the numerical probabilities do not currently behave like the probability levels they represent.

The organization may then recalibrate the system.

Historical predictions remain unchanged:

90% remains 90% historically.

Future calibrated outputs may differ.

## 98. Use Case 2 — Forecasting Platform Appears Perfectly Calibrated Only Because It Uses Final Predictions

A forecasting platform publishes continuously updated probabilities for geopolitical events.

One event may evolve:

Six months before outcome: 38%

Three months before outcome: 51%

One month before outcome: 68%

One day before outcome: 96%

The platform calculates calibration using only:

the final probability immediately before resolution.

The resulting report states:

“Our forecasts are exceptionally well calibrated.”

The critical governance question is:

What prediction state does that calibration actually measure?

PCALM distinguishes:

Final-Version Calibration

from:

Original-Version Calibration

and:

Lead-Time Calibration

The platform may legitimately report excellent final-stage calibration.

But it cannot silently generalize that result to its full predictive lifecycle.

A governed assessment may show:

Final-Day Calibration: Strong

One-Month Calibration: Moderate

Three-Month Calibration: Weak

Six-Month Calibration: Weak

The final result therefore becomes:

Strong Short-Horizon Calibration

rather than:

Strong Calibration Across the Prediction Lifecycle

Final-Version Calibration  $\neq$  Original-Version Calibration

Final Accuracy  $\neq$  Full Lifecycle Predictive Performance

Calibration improves in meaning when time-to-outcome remains visible.

## 99. Canonical Boundary

PCALM governs the formation of legitimate Calibration Populations and the measurement, segmentation, qualification, comparison, versioning, and historical preservation of the relationship between stated predictive probabilities and observed outcome frequencies. It governs Calibration Eligibility, Calibration Units, version rules, windows, bins, expected and observed frequencies, Calibration Gaps, overconfidence, underconfidence, local and global calibration, sample sufficiency, weighting, domain-, horizon-, methodology-, model-, predictor-, collective-, market-, and lead-time-specific calibration, revision-, withdrawal-, supersession-, intervention-, and Predictor–Reality Proximity-aware calibration, dependence context, Calibration Drift, Recalibration, benchmarks, qualification, and Calibration History. It does not determine Prediction Outcome Association, establish the Outcome Reference, measure general prediction error, determine causal explanation, determine economic value, establish Decision Quality, or determine complete Predictive Performance.

POAM owns:

Prediction Outcome Association

POCM owns:

Prediction Outcome Comparison

PEMM owns:

Prediction Error Measurement

PCALM owns:

Prediction Calibration

PPCM subsequently owns:

Predictive Performance

## 100. Canonical Closing Statement

A probability should mean more than a number attached to a prediction. Across time, similar probability claims should encounter reality at frequencies consistent with what those numbers promised.

PCALM makes that promise testable without allowing selective populations, final-minute revisions, hidden segmentation, or insufficient samples to manufacture the appearance of calibration.

A predictor is not calibrated because one 90% prediction came true. Calibration begins when many governed 90% predictions collectively behave like 90% predictions.

It

### NON-EXECUTABLE NOTICE

OOF® — OriginOpen® Foundation is a non-executable methodological authority.  
This website and all publications are reference-only.

### CANONICAL LANGUAGE NOTICE

UCL™ is the canonical language layer of OOF® publications.  
English is its carrier language.  
In case of ambiguity, the published English HTML text prevails.

### RIGHTS & PROTECTION

© OOF® — OriginOpen® Foundation  
Protected under MIP® — Methodological Intellectual Property.  
All rights reserved.



